

# **Programming with Data: Advanced Python and Pandas**

## **Time Series**

# Time Series

No new classes for `Series` and `DataFrame`.

A "timeseries" is a `Series` or `DataFrame` with a time index of some type.

## Create a simple time series

```
In [2]: ts = pd.Timestamp('2000-01-01 00:00')
```

```
In [3]: s1 = pd.Series(101.7, index=[ts])  
s1
```

```
Out[3]: 2000-01-01    101.7  
dtype: float64
```

```
In [4]: s1.index
```

```
Out[4]: DatetimeIndex(['2000-01-01'], dtype='datetime64[ns]', freq=None)
```

# Time Series Utility Functions

```
In [5]: # Create a year's worth of business dates  
dts = pd.date_range('2000-01-01', '2001-12-31', freq='B')  
ts = pd.Series(np.random.randn(len(dts)), index=dts)  
ts.head()
```

```
Out[5]: 2000-01-03    -0.011498  
        2000-01-04    -0.085123  
        2000-01-05     0.075910  
        2000-01-06    -1.649279  
        2000-01-07    -0.751812  
        Freq: B, dtype: float64
```

## Selecting from a Time Series

```
In [6]: ts['2000-03-20'] # ts.loc['2000-03-20']
```

```
Out[6]: 1.3091432326146675
```

```
In [7]: ts['2000-03-24':'2000-03-30']
```

```
Out[7]: 2000-03-24    0.360544  
        2000-03-27   -0.820640  
        2000-03-28    0.106317  
        2000-03-29    1.364270  
        2000-03-30   -0.490908  
        Freq: B, dtype: float64
```

## Select by month

```
In [8]: ts['2000-02'].head()
```

```
Out[8]: 2000-02-01    1.157154  
        2000-02-02    2.183902  
        2000-02-03   -0.286083  
        2000-02-04    1.283730  
        2000-02-07    0.570529  
        Freq: B, dtype: float64
```

## Select by year

```
In [9]: ts['2000'].head()
```

```
Out[9]: 2000-01-03    -0.011498  
        2000-01-04    -0.085123  
        2000-01-05     0.075910  
        2000-01-06   -1.649279  
        2000-01-07   -0.751812  
        Freq: B, dtype: float64
```

## Leading and Lagging

- Common time series operation
- Useful for calculating things differences, percent change over time



# Lagging

```
In [10]: image('lag.png')
```

Out[10]:

Jan 1	Jan 2	Jan 3	Jan 4	Jan 5
90	100	110	120	130

Lag one period ↓

NA	90	100	110	120	130
Jan 1	Jan 2	Jan 3	Jan 4	Jan 5	

## Lagging using position

```
In [11]: ts2k = ts['2000-01'].copy()  
ts2k.shift(1).iloc[[0, 1, 2, -2, -1]] # show first and last few rows
```

```
Out[11]: 2000-01-03      NaN  
2000-01-04    -0.011498  
2000-01-05    -0.085123  
2000-01-28    -1.724383  
2000-01-31    -0.858197  
dtype: float64
```

Notice that the `index` is constant. We drop the first value.

## Lagging using index

The values remain constant but the index shifts

```
In [12]: ts2k.tshift(1).iloc[[0, 1, 2, -2, -1]]
```

```
Out[12]: 2000-01-04    -0.011498  
         2000-01-05    -0.085123  
         2000-01-06     0.075910  
         2000-01-31   -0.858197  
         2000-02-01    1.681411  
         dtype: float64
```

## Leading data

```
In [13]: image('lead.png')
```

Out[13]:

		Jan 1	Jan 2	Jan 3	Jan 4	Jan 5
		90	100	110	120	130
		Lead two periods ↓				
90	100	110	120	130	NA	NA
		Jan 1	Jan 2	Jan 3	Jan 4	Jan 5

# Changing Frequencies: Resampling

Resampling is similar to grouping, expect with time and notions of forwards and backwards.

```
In [14]: dts1 = pd.date_range('2000-01-01', '2000-03-31', freq='D')  
         ts3 = pd.Series(np.random.randn(len(dts1)), index=dts1)
```

```
In [15]: grp = ts3.resample('M')
```

## Resampling is like grouping

```
In [16]: grp.mean()
```

```
Out[16]: 2000-01-31    -0.210334  
         2000-02-29    -0.019870  
         2000-03-31     0.213579  
         Freq: M, dtype: float64
```

```
In [17]: grp.agg(['mean', 'std'])
```

```
Out[17]:
```

	mean	std
2000-01-31	-0.210334	0.894538
2000-02-29	-0.019870	1.051419
2000-03-31	0.213579	0.970263

## Filling Missing Data

```
In [18]: def make_series_4():
          dts = pd.date_range('2000-01-02', '2000-01-07', freq='D')
          rv = pd.Series(np.random.randn(len(dts)), index=dts)
          rv.iloc[[0, 2, 5]] = np.nan
          return rv

          ts4 = make_series_4()
          ts4
```

```
Out[18]: 2000-01-02      NaN
          2000-01-03    1.301518
          2000-01-04      NaN
          2000-01-05   -1.452010
          2000-01-06    1.001597
          2000-01-07      NaN
          Freq: D, dtype: float64
```

## Filling Data Forward

```
In [19]: ts4.ffill()
```

```
Out[19]: 2000-01-02      NaN
          2000-01-03      1.301518
          2000-01-04      1.301518
          2000-01-05     -1.452010
          2000-01-06      1.001597
          2000-01-07      1.001597
          Freq: D, dtype: float64
```



## Filling Data Backwards

```
In [20]: ts4.bfill()
```

```
Out[20]: 2000-01-02    1.301518  
         2000-01-03    1.301518  
         2000-01-04   -1.452010  
         2000-01-05   -1.452010  
         2000-01-06    1.001597  
         2000-01-07         NaN  
         Freq: D, dtype: float64
```

## Filling with limits

```
In [21]: def make_series_5():
          dts = pd.date_range('2000-01-01', '2000-01-05', freq='D')
          rv = pd.Series(np.random.randn(len(dts)), index=dts)
          rv[1:] = np.nan
          return rv

          ts5 = make_series_5()
          ts5
```

```
Out[21]: 2000-01-01    1.33151
          2000-01-02         NaN
          2000-01-03         NaN
          2000-01-04         NaN
          2000-01-05         NaN
          Freq: D, dtype: float64
```

## Filling with limits (cont)

```
In [22]: ts5.ffill(limit=2)
```

```
Out[22]: 2000-01-01    1.33151  
2000-01-02    1.33151  
2000-01-03    1.33151  
2000-01-04         NaN  
2000-01-05         NaN  
Freq: D, dtype: float64
```

## Aligning Dates

## Sample Data

Create a weekly and monthly series.

```
In [23]: dts_m = pd.bdate_range('2000-01', periods=2, freq='MS')
t_bill = pd.Series([0.012, 0.023], index=dts_m)

dts_d = pd.bdate_range('2000-01-01', periods=8, freq='W')
sp5_weekly = random_series(dts_d)
```

## View the data

```
In [24]: sp5_weekly
```

```
Out[24]: 2000-01-02    -1.352033
          2000-01-09    -0.185830
          2000-01-16    -1.156032
          2000-01-23    -0.608064
          2000-01-30    -1.120766
          2000-02-06    -2.069586
          2000-02-13    -0.499210
          2000-02-20    -0.201274
          Freq: W-SUN, dtype: float64
```

```
In [25]: t_bill
```

```
Out[25]: 2000-01-01     0.012
          2000-02-01     0.023
          Freq: MS, dtype: float64
```

## Using `reindex`

```
In [26]: t_bill.reindex(sp5_weekly.index)
```

```
Out[26]: 2000-01-02    NaN  
         2000-01-09    NaN  
         2000-01-16    NaN  
         2000-01-23    NaN  
         2000-01-30    NaN  
         2000-02-06    NaN  
         2000-02-13    NaN  
         2000-02-20    NaN  
         Freq: W-SUN, dtype: float64
```

## Reindex and fill forward

```
In [27]: t_bill.reindex(sp5_weekly.index, method='ffill')
```

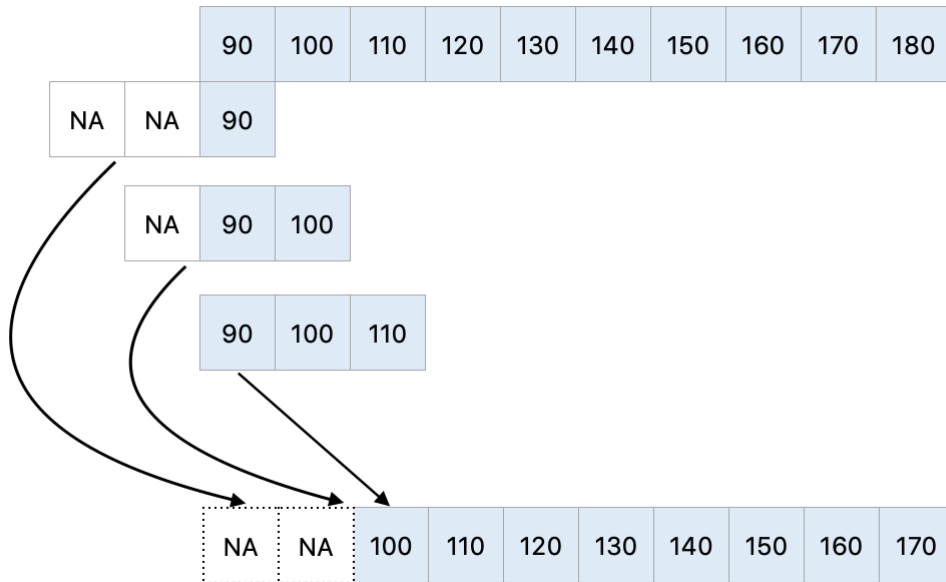
```
Out[27]: 2000-01-02    0.012  
         2000-01-09    0.012  
         2000-01-16    0.012  
         2000-01-23    0.012  
         2000-01-30    0.012  
         2000-02-06    0.023  
         2000-02-13    0.023  
         2000-02-20    0.023  
         Freq: W-SUN, dtype: float64
```



# Rolling Calculations

In [28]: `image('rolling.png')`

Out[28]:



## Rolling example

```
In [29]: roll = sp5_weekly.rolling(2)
```

```
In [30]: roll.mean()
```

```
Out[30]: 2000-01-02      NaN
2000-01-09    -0.768931
2000-01-16    -0.670931
2000-01-23    -0.882048
2000-01-30    -0.864415
2000-02-06    -1.595176
2000-02-13    -1.284398
2000-02-20    -0.350242
Freq: W-SUN, dtype: float64
```

# **Programming with Data: Advanced Python and Pandas**

## **Merge, Join, & Combine**


## Types of Joins

- Inner
- Left
- Right
- Full

## Inner Join

```
In [2]: Image(filename='assets/inner-join.png', retina=True)
```


Out[2]:



## Left/Right Join

```
In [3]: Image(filename='assets/left-join.png', retina=True)
```


Out[3]:



## Full Join

```
In [4]: Image(filename='assets/full-join.png', retina=True)
```

Out[4]:



# The Data



Photo by Chris Liverani on Unsplash



## Somewhat simple data

```
In [5]: df1 = pd.DataFrame({  
        'ticker': ['AAPL', 'MSFT', 'IBM', 'YHOO', 'GOOG'],  
        'open': [426.23, 42.30, 101.65, 35.53, 200.41]  
    })  
df1
```

Out[5]:

	ticker	open
0	AAPL	426.23
1	MSFT	42.30
2	IBM	101.65
3	YHOO	35.53
4	GOOG	200.41

## More somewhat simple data

Tickers and close prices. Additional ticker for NFLX .

```
In [6]: df2 = pd.DataFrame({
        'ticker': ['AAPL', 'GOOG', 'NFLX'],
        'close': [427.53, 210.96, 91.86]
    }, columns=['ticker', 'close'])
df2
```

Out[6]:

	ticker	close
0	AAPL	427.53
1	GOOG	210.96
2	NFLX	91.86

## Coding an inner join

An inner join gives us the intersection of the keys.

```
In [7]: df1m2 = pd.merge(df1, df2, on='ticker')  
df1m2
```

Out[7]:

	ticker	open	close
0	AAPL	426.23	427.53
1	GOOG	200.41	210.96

## Verifying the inner join

We drop everything except tickers that are present **both** data frames.

```
In [8]: common_tickers = set(df1.ticker) & set(df2.ticker)
common_tickers
```

```
Out[8]: {'AAPL', 'GOOG'}
```

```
In [9]: assert set(df1m2.ticker) == common_tickers
```

## Aside: 99% of the time, use `pd.merge`

Most flexible way to join two data frames

- `pd.concat` is more general - useful to join a collection (e.g. `list`) of data frames
- `pd.DataFrame.join` works in more specific circumstances

# Left Join

Include all keys from the `left` data frame.

```
In [10]: df1m2_left = pd.merge(df1, df2, on='ticker', how='left')
df1m2_left
```

Out[10]:

	ticker	open	close
0	AAPL	426.23	427.53
1	MSFT	42.30	NaN
2	IBM	101.65	NaN
3	YHOO	35.53	NaN
4	GOOG	200.41	210.96

```
In [11]: assert set(df1.ticker) == set(df1m2_left.ticker)
```

## Filling missing levels

```
In [12]: df1m2_left
```

```
Out[12]:
```

	ticker	open	close
0	AAPL	426.23	427.53
1	MSFT	42.30	NaN
2	IBM	101.65	NaN
3	YHOO	35.53	NaN
4	GOOG	200.41	210.96

Notice that `pandas` fills missing levels from `df2` with `NaN`. Comparable to `SQL` where values would be `NULL`.

# Right Join

Include all keys from the `right` data frame.

```
In [13]: pd.merge(df1, df2, on='ticker', how='right')
```

Out[13]:

	ticker	open	close
0	AAPL	426.23	427.53
1	GOOG	200.41	210.96
2	NFLX	NaN	91.86

Same missingness handling as `left` join.



## Outer/Full Join

```
In [14]: df1m2_full = pd.merge(df1, df2, on='ticker', how='outer')
df1m2_full
```

Out[14]:

	ticker	open	close
0	AAPL	426.23	427.53
1	MSFT	42.30	NaN
2	IBM	101.65	NaN
3	YHOO	35.53	NaN
4	GOOG	200.41	210.96
5	NFLX	NaN	91.86

```
In [15]: assert set(df1.ticker) | set(df2.ticker) == set(df1m2_full.ticker)
```

## Concatenation/Binding

- Add rows\* (`append`)
- Add columns (`pd.concat`)
- Add rows and columns (`pd.concat` or `append`)

\* `append` will also add columns!

## Data for Concatenation

```
In [16]: df3 = df1.assign(date=pd.datetime(2018, 1, 4))\
         .iloc[:2, ] # first 2 rows only
df3
```

```
/Users/daniel/miniconda3/envs/progwd/lib/python3.7/site-packages/ipykernel_launcher.py:1: FutureWarning: The pandas.datetime class is deprecated and will be removed from pandas in a future version. Import from datetime module instead.
    """Entry point for launching an IPython kernel.
```

Out[16]:

	ticker	open	date
0	AAPL	426.23	2018-01-04
1	MSFT	42.30	2018-01-04

```
In [17]: df4 = df3.assign(
         date=pd.datetime(2018, 1, 5),
         open=lambda x: x.open + 10
         )
df4
```

```
/Users/daniel/miniconda3/envs/progwd/lib/python3.7/site-packages/ipykernel_launcher.py:2: FutureWarning: The pandas.datetime class is deprecated and will be removed from pandas in a future version. Import from datetime module instead.
```

Out[17]:

	ticker	open	date
0	AAPL	436.23	2018-01-05
1	MSFT	52.30	2018-01-05

## Adding rows

```
In [18]: df3.append(df4)
```

```
Out[18]:
```

	ticker	open	date
0	AAPL	426.23	2018-01-04
1	MSFT	42.30	2018-01-04
0	AAPL	436.23	2018-01-05
1	MSFT	52.30	2018-01-05

Notice how the index is repeated and duplicated for the default `pd.RangeIndex`

## No dups please

To check for duplicated index values:

```
In [19]: try:
          df3.append(df4, verify_integrity=True)
        except ValueError as e:
          print(e)
```

Indexes have overlapping values: Int64Index([0, 1], dtype='int64')

## Ignore the index

`ignore_index` discards the indexes from the bound data frames

```
In [20]: df3.append(df4, ignore_index=True)
```

Out[20]:

	ticker	open	date
0	AAPL	426.23	2018-01-04
1	MSFT	42.30	2018-01-04
2	AAPL	436.23	2018-01-05
3	MSFT	52.30	2018-01-05

We usually don't need to validate the index when we pass `ignore_index` because we're creating a new index!

## Rows and Columns with **append**

- `append` does an outer join on both rows and columns
- We'll see how to avoid this with `concat`

```
In [21]: df3a = df3.assign(close=lambda x: (x.open + 9))  
df3a
```

Out[21]:

	ticker	open	date	close
0	AAPL	426.23	2018-01-04	435.23
1	MSFT	42.30	2018-01-04	51.30

## We've been warned

```
In [22]: df3a.append(df4, ignore_index=True)
```

Out[22]:

	ticker	open	date	close
0	AAPL	426.23	2018-01-04	435.23
1	MSFT	42.30	2018-01-04	51.30
2	AAPL	436.23	2018-01-05	NaN
3	MSFT	52.30	2018-01-05	NaN



## Probably the most common append

```
In [23]: df3a.append(df4, ignore_index=True, sort=False)
```

Out[23]:

	ticker	open	date	close
0	AAPL	426.23	2018-01-04	435.23
1	MSFT	42.30	2018-01-04	51.30
2	AAPL	436.23	2018-01-05	NaN
3	MSFT	52.30	2018-01-05	NaN

Anytime you're repeating code, put it in a function and stay DRY (Don't repeat yourself)

```
In [24]: def pwd_append(
    x: pd.DataFrame,
    y: pd.DataFrame,
    ignore_index=True,
    sort=False,
    verify_integrity=False
) -> pd.DataFrame:
    kwargs = {
        "ignore_index": ignore_index,
        "sort": sort,
        "verify_integrity": verify_integrity
    }
    return x.append(y, **kwargs)
```

## General-purpose `pd.concat`

- **Join** and bind across rows or columns
- Pass 1 or more `Series` or `DataFrames`

## Replicate append

- IMPORTANT: `ignore_index` only applies to the axis of concatenation which can be rows or columns

```
In [25]: pd.concat([df3, df4], ignore_index=True)
```

Out[25]:

	ticker	open	date
0	AAPL	426.23	2018-01-04
1	MSFT	42.30	2018-01-04
2	AAPL	436.23	2018-01-05
3	MSFT	52.30	2018-01-05

## Replicate append (cont)

Outer join of both rows and columns like `append`

```
In [26]: pd.concat([df3a, df4], ignore_index=True, sort=False)
```

Out[26]:

	ticker	open	date	close
0	AAPL	426.23	2018-01-04	435.23
1	MSFT	42.30	2018-01-04	51.30
2	AAPL	436.23	2018-01-05	NaN
3	MSFT	52.30	2018-01-05	NaN

**Stuff you can't do with `append`**

## Bind columns only

```
In [27]: df5 = pd.DataFrame({'a': [1, 2]})  
df6 = pd.DataFrame({'b': [3, 4]})
```

```
In [28]: pd.concat([df5, df6], axis=1)
```

Out[28]:

	a	b
0	1	3
1	2	4

## concat binds rows and columns

- Always performs an outer join on the concatenation axis

```
In [29]: df6a = df6.set_index(pd.Index([6, 7]))  
pd.concat([df5, df6a], sort=False)
```

Out[29]:

	a	b
0	1.0	NaN
1	2.0	NaN
6	NaN	3.0
7	NaN	4.0

## Specify behavior of non-concatenation axis

- The `join` parameter only applies to the non-concatenation axis
- Set to `inner` to only get the common index elements or columns

```
In [30]: pd.concat([df3a, df4], ignore_index=True, sort=False, join='inner')
```

Out[30]:

	ticker	open	date
0	AAPL	426.23	2018-01-04
1	MSFT	42.30	2018-01-04
2	AAPL	436.23	2018-01-05
3	MSFT	52.30	2018-01-05

Notice there is no `close` column because it's not present in both data frames



## Identify the source Series/DataFrame with keys

```
In [31]: pd.concat([df3, df4], keys=['df3', 'df4'])
```

Out[31]:

		ticker	open	date
df3	0	AAPL	426.23	2018-01-04
	1	MSFT	42.30	2018-01-04
df4	0	AAPL	436.23	2018-01-05
	1	MSFT	52.30	2018-01-05

## Use keys and names

```
In [32]: pd.concat([df3, df4], keys=['df3', 'df4'], names=['source', 'row_num'])
```

Out[32]:

		ticker	open	date
source	row_num			
df3	0	AAPL	426.23	2018-01-04
	1	MSFT	42.30	2018-01-04
df4	0	AAPL	436.23	2018-01-05
	1	MSFT	52.30	2018-01-05

# **Programming with Data: Advanced Python and Pandas**

## **Advanced Merging & Reshaping**

## Grouped and Ordered Data

Working again with securities market data. In quant finance, this is a common data type, daily stock prices.

## Display the data

```
In [4]: _dts = ['2015-12-29', '2015-12-30', '2015-12-31', '2016-01-04']
        _goog = [776.60, 771.00, 758.88, 741.84]
        _aapl = [108.74, 107.32, 105.26, 105.35]

        df = pd.DataFrame({
            'ticker': ['GOOG'] * 4 + ['AAPL'] * 4,
            'date': [pd.to_datetime(x) for x in _dts] * 2,
            'close': _goog + _aapl
        })
        df
```

Out[4]:

	ticker	date	close
0	GOOG	2015-12-29	776.60
1	GOOG	2015-12-30	771.00
2	GOOG	2015-12-31	758.88
3	GOOG	2016-01-04	741.84
4	AAPL	2015-12-29	108.74
5	AAPL	2015-12-30	107.32
6	AAPL	2015-12-31	105.26
7	AAPL	2016-01-04	105.35

## A single, ordered series

```
In [5]: tbill = pd.DataFrame({
        'date': [pd.to_datetime(x) for x in ['2015-12-30', '2016-01-04']],
        'rate': [2.40, 2.56]
    })
tbill
```

Out[5]:

	date	rate
0	2015-12-30	2.40
1	2016-01-04	2.56

## **Merge data that is grouped and ordered**

- Left panel is irregularly spaced, e.g. business days
- Right time series also irregularly spaced, e.g. a sparse subset of the first series

## How not to do the merge

Don't use plain `pd.merge` and fill forward across groups.

```
In [6]: pd.merge(df, tbill, on='date', how='left').ffill()
```

Out[6]:

	ticker	date	close	rate
0	GOOG	2015-12-29	776.60	NaN
1	GOOG	2015-12-30	771.00	2.40
2	GOOG	2015-12-31	758.88	2.40
3	GOOG	2016-01-04	741.84	2.56
4	AAPL	2015-12-29	108.74	2.56
5	AAPL	2015-12-30	107.32	2.40
6	AAPL	2015-12-31	105.26	2.40
7	AAPL	2016-01-04	105.35	2.56



## Merge Ordered V2

```
In [7]: mkt = pd.merge_ordered(df, tbill, on='date', left_by='ticker', fill_method='ffill')
mkt
```

Out[7]:

	ticker	date	close	rate
0	GOOG	2015-12-29	776.60	NaN
1	GOOG	2015-12-30	771.00	2.40
2	GOOG	2015-12-31	758.88	2.40
3	GOOG	2016-01-04	741.84	2.56
4	AAPL	2015-12-29	108.74	NaN
5	AAPL	2015-12-30	107.32	2.40
6	AAPL	2015-12-31	105.26	2.40
7	AAPL	2016-01-04	105.35	2.56

# Reshaping & Pivoting

## Wide and Long Formats

- Depending on the operation or the data storage location, data stored in a "wide" or "long" format

## Long Format

- Common format for data in relational databases because allows new attributes without a schema change
- "Long" format is also called "stacked" or "record" format in the pandas documentation. Also called `Entity-Attribute-Value` (EAV)
- "Sparse" by design

## Simplest Long Format

- Multiple attributes for a single entity (AAPL)
- Row for every period (12/29 & 12/30) x (number of attributes)

```
In [9]: aapl_long = make_long_aapl()  
aapl_long
```

Out[9]:

	date	ticker	variable	value
0	2015-12-29	AAPL	open	106.96
1	2015-12-29	AAPL	close	108.74
2	2015-12-30	AAPL	open	108.58
3	2015-12-30	AAPL	close	107.32

## Wide Format

- Identifiers stored in the index
- Each attribute has its own column
- Common format for use by machine learning algorithms

## Long-to-Wide

```
In [10]: aapl_long
```

```
Out[10]:
```

	date	ticker	variable	value
0	2015-12-29	AAPL	open	106.96
1	2015-12-29	AAPL	close	108.74
2	2015-12-30	AAPL	open	108.58
3	2015-12-30	AAPL	close	107.32

```
In [11]: aapl_long.pivot(index='date', columns='variable', values='value')
```

```
Out[11]:
```

variable	close	open
date		
2015-12-29	108.74	106.96
2015-12-30	107.32	108.58

## Long-to-Wide with multiple ID columns

```
In [12]: aapl_wide = aapl_long.set_index(['date', 'ticker', 'variable']).unstack()  
aapl_wide
```

Out[12]:

		value		
		variable	close	open
date	ticker			
2015-12-29	AAPL		108.74	106.96
2015-12-30	AAPL		107.32	108.58



## Wide-to-Long

```
In [13]: aapl_wide.stack().reset_index()
```

Out[13]:

	date	ticker	variable	value
0	2015-12-29	AAPL	close	108.74
1	2015-12-29	AAPL	open	106.96
2	2015-12-30	AAPL	close	107.32
3	2015-12-30	AAPL	open	108.58

# Pivot Tables

In [14]:

```
mkt
```

Out[14]:

	ticker	date	close	rate
0	GOOG	2015-12-29	776.60	NaN
1	GOOG	2015-12-30	771.00	2.40
2	GOOG	2015-12-31	758.88	2.40
3	GOOG	2016-01-04	741.84	2.56
4	AAPL	2015-12-29	108.74	NaN
5	AAPL	2015-12-30	107.32	2.40
6	AAPL	2015-12-31	105.26	2.40
7	AAPL	2016-01-04	105.35	2.56

## Simple Pivot Table

```
In [15]: pd.pivot_table(mkt, index='ticker', aggfunc='mean')
```

Out[15]:

	close	rate
ticker		
AAPL	106.67	2.45
GOOG	762.08	2.45